

Given an array of integers, and x , find how many pairs of elements of the array sum to x , i.e., how many indexes $i \neq j$ are there with $arr[i] + arr[j] = x$? Array is not sorted, and may contain duplicate elements. Solve the problem using balanced BST with $RT = O(n \log n)$. For example, If $arr = \{3, 3, 4, 5, 3, 5, 4\}$ then $howMany(A, 8)$ returns 7.

Answer #1

```
static int howMany(int[] arr, int x) {
    Map<Integer, Integer> map = new TreeMap<>();
    for (int e : arr) {
        Integer c = map.get(e);
        if (c == null) map.put(e, 1);
        else map.put(e, c + 1);
    }
    int count = 0;
    for (Map.Entry<Integer, Integer> me : map.entrySet()) {
        int e = me.getKey(); int c = me.getValue();
        if (e * 2 > x) break;
        int c2 = map.get(x - e);
        if (c2 == null) continue;
        else { if (e == x - e) { count += c * (c - 1) / 2; }
              else { if (2 * e <= x) count += c * c2; }
            }
    }
    return count;
}
```

slight answer:

```
Map<Integer, Integer> map = new TreeMap<>();
for (e : arr) { map.put(e, 0); map.put(x - e, 0); }
count = 0;
for (e : arr) {
    count += map.get(x - e);
    map.put(e, map.get(e) + 1);
}
return count;
```

Given an array, return those elements that occur exactly once, in the same order in which they appear in the given array. Solve the problem using balanced BST, with $RT = O(n \log n)$.

```
T extends Comparable<? super T>
static<T> Object[] exactlyOnce(T[] arr) {
    TreeMap<T, Integer> map = new TreeMap<>();
    int unique = 0;
    for (T e : arr) {
        Integer c = map.get(e);
        if (c == null) {map.put(e, 1); unique++;}
        else {map.put(e, c+1); if (c == 1) unique--;}
    }
    // unique = number of unique elements
    Object[] result = new Object[unique];
    int i = 0;
    for (T e : arr) {
        Integer c = map.get(e);
        if (c == 1) {result[i++] = e;}
    }
    return result;
}
```

Given an array of integers, find the length of a longest streak of consecutive integers that occur in it (not necessarily contiguously). For example, if `arr = {1,7,9,4,1,7,4,8,7,1}`. `longestStreak(arr)` returns 3, corresponding to the streak {7,8,9} of consecutive integers that occur in the array. Solve using balanced BST, with RT = $O(n \log n)$.

```
static int longestStreak(int[] arr) {
```

```
    TreeSet<Integer> set = new TreeSet<>();
    for (int e : arr) set.add(e);
    prev = set.first();
    max = 1; // just prev
    current = 0;
    // LI: max = Length of max streak seen so far, current = length of streak upto current elem
    for (int e : set) { // in sorted order for TreeSet
        if (prev + 1 == e) { current++; }
        else { current = 1; }
        prev = e;
        if (current > max) max = current;
    }
    return max;
}
```

Build height-balanced BST from sorted array.

```
treeFromArray( arr ):
```

```
    tree ← treeFromArray( arr, 0, arr.length )
```

```
    return new BST( tree, arr.length )
```

```
// Build a balanced tree from arr[ i..i+n-1 ]
```

```
treeFromArray( arr, i, n ): // RT = O(n).
```

```
    if n <= 0 then
```

```
        return null
```

```
    else
```

```
        Ln ← n / 2    // arr[ i..i+Ln-1 ]
```

```
        Rn ← n - Ln - 1 // arr[ i+Ln+1..i+n-1 ]
```

```
        root ← i + Ln
```

```
        left ← treeFromArray ( arr, i, Ln )
```

```
        right ← treeFromArray ( arr, root+1, Rn )
```

```
        return new Tree( arr[ root ], left, right )
```

Let the running time of `treeFromArray(arr, i, n)`

be $T(n)$. The recurrence for the running time is:

$T(0) = 1$ (base case of the algorithm)

$T(n) = T(Ln) + T(Rn) + c$, where c is a constant

with the running time of the other lines of the

program, other than the recursive calls.

When $n+1$ is an exact power of 2, the array splits equally, and a complete binary tree is returned.

For this case, the recurrence is $T(n) = 2T(n/2) + c$.

In class we used the Master method to show that the solution to this recurrence is $T(n) = O(n)$.

Hashing: subset of dictionary operations: add, contains, remove.

A function h , known as hash function, maps elements to non-negative integers in $[0, n-1]$ where the table size is chosen to be n . Then x will be placed in $\text{table}[h(x)]$, if possible.

Design goals:

1. Choose n proportional to number of elements in dictionary: $\lambda = \text{size} / n$, the load factor, is $O(1)$.
2. For any two keys x and y , $\Pr\{h(x) = h(y)\} = 1/n$.
3. Pseudorandom function: $h(1), h(2), h(3), \dots$ should be indistinguishable from a random sequence.
4. Deterministic, and easy to compute.

Implementation sketch:

add(x): Place x in $\text{table}[h(x)]$	contains(x): Is x in $\text{table}[h(x)]$?	remove(x): remove x from $\text{table}[h(x)]$
---	---	---

Collision resolution: What do you do if $\text{add}(x)$ finds $\text{table}[h(x)]$ is already occupied by another element?

- (1) Separate chaining (known as open hashing): each entry of the hash table is a linked list of elements.
- (2) Open addressing (closed hashing): each entry of the hash table can store only one element (or a small, fixed number of elements). Many schemes are available for collision resolution.

Java: Hash tables use separate chaining. Hash function is called `hashCode()`, and $h(x)$ is a function of $x.\text{hashCode}()$ and n . Table size is automatically adjusted based on load factor, and system tries to keep the load factor to be less than 0.5. In the base class of the object hierarchy, `Object`, `hashCode` is defined to be the address of the object. This is not a good hash function. Wrapper classes override it. User-defined classes that need to be used as keys in hashing should implement `hashCode()` and `equals()` methods.

The lengths of Java's hash tables are powers of 2 to simplify calculations. Bit operations are used to mangle the integer given by `hashCode()` to avoid problems created by poorly defined hash functions.

```
// Code extracted from Java's HashMap:
static int hash(int h) {
    // This function ensures that hashCodes that differ only by
    // constant multiples at each bit position have a bounded
    // number of collisions (approximately 8 at default load factor).
    h ^= (h >>> 20) ^ (h >>> 12);
    return h ^ (h >>> 7) ^ (h >>> 4);
}
static int indexFor(int h, int length) {    // length = table.length is a power of 2
    return h & (length-1);
}
// Key x is stored at table[ hash( x.hashCode() ) & ( table.length - 1 ) ].
```

Java hash tables: `HashSet`, `HashMap`, `LinkedHashSet`, `ConcurrentHashMap`, `HashTable`.

HashSet: implementation of `Set` interface. Main operations: `add`, `contains`, `remove`, `iterator`. `add(x)` is rejected if x is already in the set. `HashSet` is implemented using `HashMap`.

HashMap: Implementation of `Map` interface (key/value pairs). Main ops: `get`, `put`, `containsKey`, `remove`, `iterator`. `put` operation replaces value if key already exists in map. `get` returns null if key does not exist.

LinkedHashSet: like `HashSet`, but `iterator` goes through elements in order of `add`.

ConcurrentHashMap, HashTable: synchronized, suitable for multi-threaded applications.