


Basis of all shortest path algorithms: subpath of a shortest path is a shortest path. In other words, if a shortest path from s to v is composed of a path from s to u and the edge (u,v) , then $\delta(s,v) = \delta(s,u) + w(u,v)$.



Therefore, shortest paths can be encoded as an up-tree, where each node stores its predecessor in a shortest path from s to that node. In the example above, we can set $v.\pi = u$. The following utility functions are used by all shortest path algorithms with edge weights:

initialize(s): for $u \in V$ do $u.d \leftarrow \infty$ $u.\pi \leftarrow \text{null}$ $u.\text{seen} \leftarrow \text{false}$ $s.d \leftarrow 0$	boolean relax(u, v, e): if $v.d > u.d + e.\text{weight}$ then $v.d \leftarrow u.d + e.\text{weight}$ $v.\pi \leftarrow u$ return true return false
---	--

DAG-shortest-paths algorithm: In a DAG, the nodes in any path are in strictly increasing order of their topological numbers, in any topological ordering of V . This can be exploited to design the following efficient algorithm for shortest paths in DAGs:

// Pull algorithm dagSP(g, s): Find a topological ordering of g initialize(s) for $u \in V$ in topological order do // LI: for predecessors p of u , $p.d = \delta(s, p)$ for edge $e = (p, u)$ into u do relax(p, u, e)	// Push algorithm dagSP(g, s): Find a topological ordering of g initialize(s) for $u \in V$ in topological order do // LI: $u.d = \delta(s, u)$ for edge $e = (u, v)$ out of u do relax(u, v, e)
---	---

Dijkstra's algorithm: applicable in graphs without any edges of negative weight.

Idea:

- * Maintain a set of nodes S for which shortest paths are known.
- * For $v \in V - S$, store in $v.d$, the length of a shortest path from s to v that goes through only nodes of S .
- * In each iteration, select a node u in $V - S$ with minimum $u.d$, and add it to S .
- * Relax edges out of u to update distance estimates of other nodes in $V - S$.

Code resembles Prim's algorithm that uses indexed priority queues:


dijkstraSP(G, s): // Implementation using indexed priority queue of vertices // $v \in V - S$ stores in $v.d$, the weight of a shortest path from s to v that goes through only nodes of S initialize(s) // for $u \in V$ do { $u.d \leftarrow \infty$; $u.\pi \leftarrow \text{null}$; $u.\text{seen} \leftarrow \text{false}$ } $s.d \leftarrow 0$ $q \leftarrow$ new indexed priority queue of vertices with $u.d$ as priority of u while q is not empty do $u \leftarrow q.\text{remove}()$ // LI: $u.d = \delta(s, u)$ $u.\text{seen} \leftarrow \text{true}$ for Edge $e = (u, v)$ incident on u do if not $v.\text{seen}$ and relax(u, v, e) then $q.\text{decreaseKey}(v)$

Dijkstra's algorithm for shortest paths

Input: $G = (V, E)$ (directed or undirected),
 $w: E \rightarrow \mathbb{Z}^+$ ~ nonnegative weights
on edges.
Source $s \in V$ (\mathbb{R}^+)

Output: For all $u \in V$, $\delta(s, u)$, $\pi[u]$
↑
shortest path weight
from s to u
↑
predecessor of
 u in this
shortest path

Example:




$u.d$ = estimate of how far u is from s .

Execution of Dijkstra's algorithm on this example:

Order in which nodes are added to $S = \{s, D, B, E, A, F, C\}$

shortest path tree



Vertex	d	π
S	0	-
A	5	D
B	4	D
C	11 9	A
D	3	S
E	5	D
F	8 6	E

Final Exam: 8:00-10:30 AM on Mon, Dec 17.

Topics: All topics discussed in class/assignments/projects

4 cheat sheets allowed, $8\frac{1}{2}'' \times 11''$ paper - both sides
(8 pages in all)

Exam will have 3 sections.

A⁺/A/A⁻ grades can be earned only by answering

Sections 2 and 3. To get B⁺... F grades,

Section 3 need not be answered.

Idea: Maintain a set S of nodes, such that
for $u \in S$, $u.d = \delta(s, u)$.

for $v \in V - S$, $v.d = \text{length of a}$
shortest path from s to v , all of
whose internal nodes are in S ..

Nodes are kept in a priority queue with
of $V-S$ $u.d$ as priority of u .

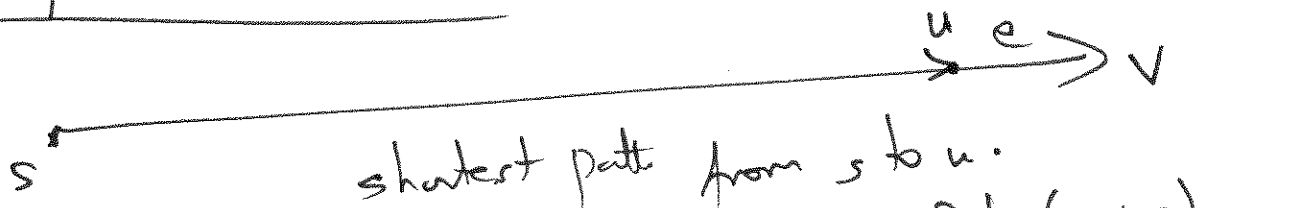
In each iteration, remove node with smallest $u.d$,
process edges out of u .

RT of Dijkstra's algorithm:

$O(|E| \log |V|)$ with indexed binary heaps

$O(|E| + |V| \log |V|)$ with Fibonacci
heaps

shortest paths in DAGs: $\delta(s, v) = \delta(s, u) + w(u, v)$.


Suppose, we know that $u.d = \delta(s, u) \rightarrow \text{Relax}(u, v, e)$
: $v.d = \delta(s, v)$

Which one of these nodes is predecessor of v in $\delta(s, v)$?




What if we wait until all predecessors of v have been processed?

Then we relax all edges into v - one of these edges is the right one.

Process nodes in topological order. $RT = O(|E| + |V|)$

Example:



Topological order = $\{v, s, t, x, y, z\}$ v as source

s as source \rightarrow

	d	π
✓ v	∞	-
✓ s	0	-
✓ t	2	s
✓ x	6	s
y	5	t x
z	3	t y

	d	π
✓ v	0	-
✓ s	5	v
✓ t	3	v
✓ x	10	s t
✓ y	7	t
✓ z	5	t

Bellman-Ford Algorithm

Input: Graph $G = (V, E)$ - directed graph
 $W: E \rightarrow \mathbb{Z}$ - positive and negative weights
 (\mathbb{R})
Source $s \in V$

Output: either: (1) For each $u \in V: \delta(s, u), \pi(u)$
if G has no negative cycles
or (2) Discover a negative cycle in G .

Negative cycle: cycle $C: \sum_{e \in C} w(e) < 0$.

Idea: Def: $d_k(u) =$ length of a shortest ^(simple) path from s to u using at most k edges.

(2) If a graph has no negative cycles,
a path P from s to u that is not simple
can be used to get a path P' from s to u
that is simple,
 $w(P') \leq w(P)$.

\Rightarrow shortest path with no constraint on simplicity
is fine.

